

Please upload your homework to Gradescope by April 12, 4:00 pm.

Please submit a single PDF directly on Gradescope

You may type your homework or scan your handwritten version. Make sure all the work is discernible.

1. In class, we learned that the solution of the least square problem satisfies the normal equation $X^T X w = X^T y$, where $X \in \mathbb{R}^{N \times M}$ and $y \in \mathbb{R}^N$. For this normal equation to have a unique solution, $X^T X$ has to be nonsingular. Prove the following statement about the necessary and sufficient conditions of the nonsingularity of $X^T X$:

The Gram matrix $X^T X$ is nonsingular if and only if X has linearly independent columns.

Hint: A matrix A has linearly independent columns if $Ax = 0$ only for $x = 0$. You may want to use the following properties of (non)singular matrix. The equation $Ax = 0$ has only the trivial solution $x = 0$ if and only if A is nonsingular. If A is singular, there exist $z \neq 0$ such that $Az = 0$.

In our set-up, with $N \gg M$, there will likely be M linearly independent vectors x_n . Therefore, the Gram matrix $X^T X$ is likely to be invertible.

2. Consider the hat matrix $H = X(X^T X)^{-1} X^T$, where $X \in \mathbb{R}^{N \times M}$ and $X^T X$ is invertible. We encountered this hat matrix when we solved the linear least squares regression problem. The vector $w = (X^T X)^{-1} X^T y$ led to the prediction $\hat{y} = Xw = X(X^T X)^{-1} X^T y = Hy$. In other words, this hat matrix H describes the linear transformation between the target y (true label) and the estimation \hat{y} . We will prove several properties of H in this question.

(a) Show that H is symmetric.

(b) Show that $H^K = H$ for any positive integer K .

(c) If I is the identity matrix of size N , show that $(I - H)^K = I - H$ for any positive integer K .

(d) Show that $\text{Trace}(H) = M$, where the trace is the sum of diagonal elements.
Hint: $\text{Trace}(AB) = \text{Trace}(BA)$.

3. Consider a linear regression problem in which we want to “weigh” different training instances differently because some of these instances are more important than others. Specifically, suppose we want to minimize

$$J(w_0, w_1) = \sum_{n=1}^N \alpha_n (w_0 + w_1 x_n - y_n)^2. \quad (1)$$

Here $\alpha_n > 0$; w_0 and w_1 are weights; x_n 's and y_n 's are scalars. In class we worked out what happens for the case where all weights (the α_n 's) are the same. In this problem, we will generalize some of those ideas to the weighted setting.

- (a) Calculate the gradient by computing the partial derivative of J with respect to each of the parameters w_0 and w_1 .
 - (b) Comment on how the α_n 's affect the linear regression problem. For example, what happens if $\alpha_i = 0$ for some i , $1 \leq i \leq N$?
 - (c) Qualitatively describe what will happen when performing gradient descent if one α_j is much greater than all other $\alpha_{j'}$'s, $j' \neq j$, $1 \leq j, j' \leq N$.
4. Consider the following objective function:

$$J(w) = \frac{1}{2} \|w\|^2 + C \left[\frac{1}{n} \sum_{i=1}^n \max(0, 1 - y_i(w^T x_i)) \right],$$

where C is a constant, $x_i \in \mathbb{R}^M$ and $y_i \in \{-1, +1\}$. We will encounter a similar loss function when we study support vector machine (SVM). Derive the stochastic gradient descent (SGD) rule to minimize this loss function.

5. In this exercise, you will implement approaches to solve the linear and polynomial regression problems. Our data consists of inputs $x_n \in \mathbb{R}$ and outputs $y_n \in \mathbb{R}$, $n \in \{1, \dots, N\}$, which are related through a target function $f(x)$. Your goal is to learn a predictor $h_w(x)$ that best approximates $f(x)$.
- (a) **Visualization** We provide two sets of data, the training data and the testing data in the two files, *regression.train.csv* and *regression.test.csv*. In each file, the first column is the input and the second column is the output. In MATLAB (or python), visualize the training data by plotting the input vs. the output. What do you observe? For example, can you make an educated guess on the effectiveness of linear regression in predicting the data?
 - (b) **Linear Regression: closed-form solution** Let us start by considering a simple linear regression model:

$$h_w(x) = w^T x = w_0 + w_1 x.$$

Recall that linear regression attempts to minimize the objective function

$$J(w) = \sum_{n=1}^N (h_w(x_n) - y_n)^2 = \|Xw - y\|^2,$$

where

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}, X = \begin{bmatrix} 1, x_1 \\ 1, x_2 \\ \vdots \\ 1, x_N \end{bmatrix}, w = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix}.$$

Note that to take into account the intercept term w_0 , we can add an additional “feature” to each instance and set it to one, e.g., $x_{i,0} = 1$. This step is equivalent to adding an additional first column to X and setting it to all ones.

In class we learned that the closed-form solution to linear regression is:

$$w^* = (X^T X)^{-1} X^T y.$$

Implement this closed-form solution in MATLAB (or python) using the training data and report w^* and $J(w)$. Also generate a plot depicting your training data and the fitted line.

- (c) **Linear Regression: gradient descent** Another way to solve linear regression is through gradient descent (GD). In gradient descent, each iteration performs the following update:

$$w_j := w_j - \eta \sum_{n=1}^N (h_w(x_n) - y_n) x_{n,j} \quad (\text{simultaneously update } w_j \text{ for all } j).$$

With each iteration of gradient descent, we expect our updated parameters w to come closer to the optimal w^* and therefore achieve a lower value of $J(w)$.

Implement the gradient descent algorithm in MATLAB (or python) using all of the following specifications for the gradient descent algorithm:

- Initialize w to be the all 0’s vector.
- Run the algorithm for 10000 iterations.
- Terminate the algorithm earlier if the value of $J(w)$ is unchanged across consecutive iterations. In this exercise, we say $J(w)$ is unchanged if $|J(w)_{t-1} - J(w)_t| < 0.0001$.
- Use a fixed learning rate η .

For different $\eta = 0.05, 0.001, 0.0001, 0.00001$, report the final $J(w)$ and the number of iterations until convergence (the number will be 10000 if the algorithm did not converge within 10000 iterations). Discuss how does the learning rate η affect the GD algorithm.

- (d) **Gradient Descent Visualization** Repeat the exercise in (c) with only $\eta = 0.05$ for only 40 iterations (updates). Initialize w to be the all 0s vector and treat initialization as iteration 0. Report w , $J(w)$ and plot the fitted line (along with the data) for iteration 0, 10, 20, 30 and 40. Compare your fitted lines and $J(w)$ (s) to what you get in (b). What do you observe over different iterations?
- (e) **Polynomial Regression** Next let us consider the more complicated case of polynomial regression, where our hypothesis is

$$h_w(x) = w^T \phi(x) = w_0 + w_1 x + w_2 x^2 + \dots + w_m x^m.$$

Note that the function is linear in the vector w of coefficients. We can therefore extend the result of linear regression by replacing the input matrix X with

$$\Phi = \begin{bmatrix} \phi(x_1)^T \\ \phi(x_2)^T \\ \vdots \\ \phi(x_N)^T \end{bmatrix}$$

where $\phi(x)$ is a vector such that $\phi_j(x) = x^j$ for $j = 0, \dots, m$.

For $m = 0, \dots, 10$, use the closed-form solution to determine the best-fit polynomial regression model on the training data.

- i. With this model, calculate the RMSE (Root-Mean-Square error), i.e., $E_{RMS} = \sqrt{J(w)/N}$, on both the training data and the test data.
- ii. Generate a plot depicting how RMSE (both training and testing) varies with model complexity (polynomial degree m).
- iii. Which degree of the polynomial would you say best fits the data? Was there evidence of under/over-fitting the training data? Use your plot to justify your answers.